


envoy
by **lyft**

Lyft's Envoy: Experiences Operating a Large Service Mesh
SREcon17
Matt Klein / @mattklein123, Software Engineer @Lyft

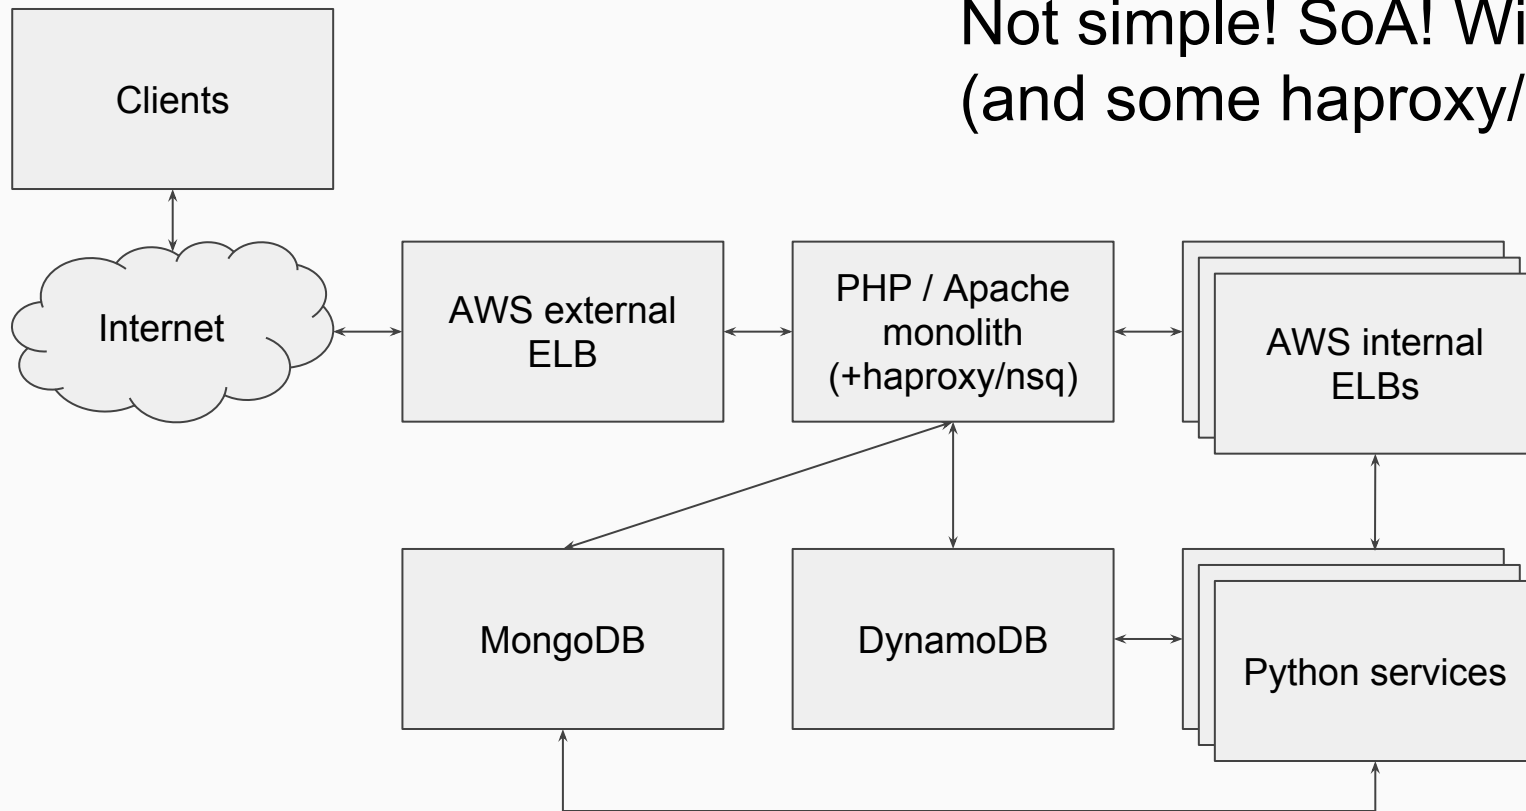


Lyft ~4 years ago



Simple! No SoA! (*but still not that simple*)

Not simple! SoA! With monolith!
(and some haproxy/nsq)



State of SoA networking in industry

- **Languages** and frameworks.
- **Protocols** (HTTP/1, HTTP/2, gRPC, databases, caching, etc.).
- **Infrastructures** (IaaS, CaaS, on premise, etc.).
- Intermediate **load balancers** (AWS ELB, F5, etc.).
- **Observability** output (stats, tracing, and logging).
- Implementations (often partial) of **retry**, **circuit breaking**, **rate limiting**, **timeouts**, and other distributed systems best practices.
- **Authentication** and **Authorization**.
- Per language **libraries** for service calls.

State of SoA networking in industry



A really big and confusing mess...

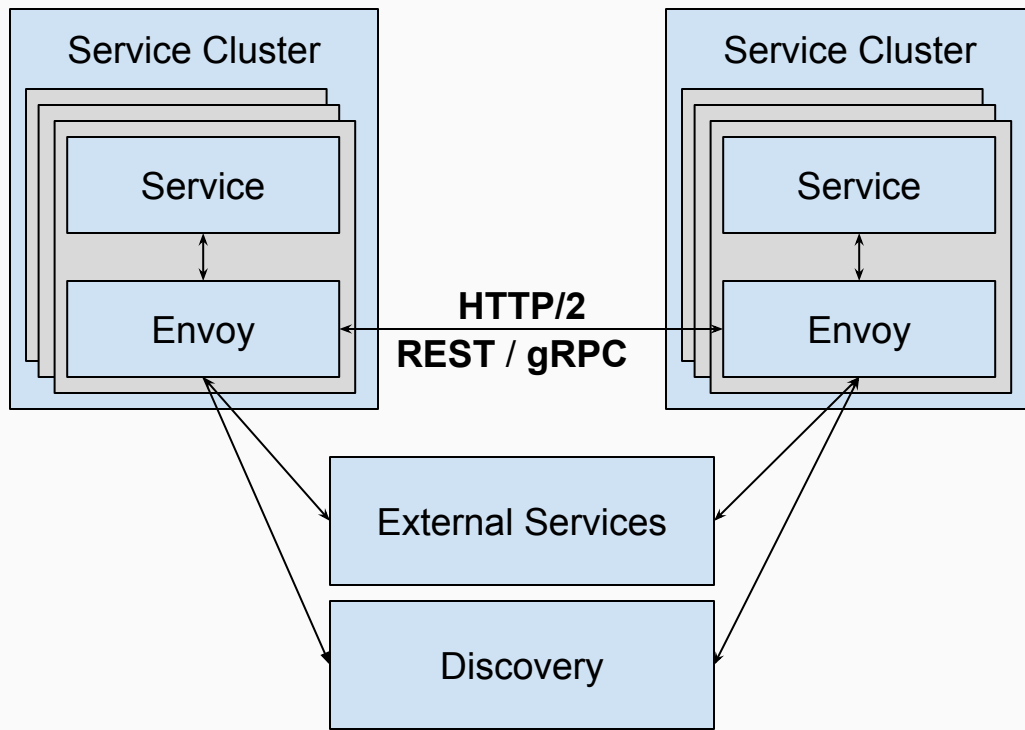
The network should be transparent to applications. When network and application problems do occur it should be easy to determine the source of the problem.

This sounds great! But it turns out it's really, really hard.

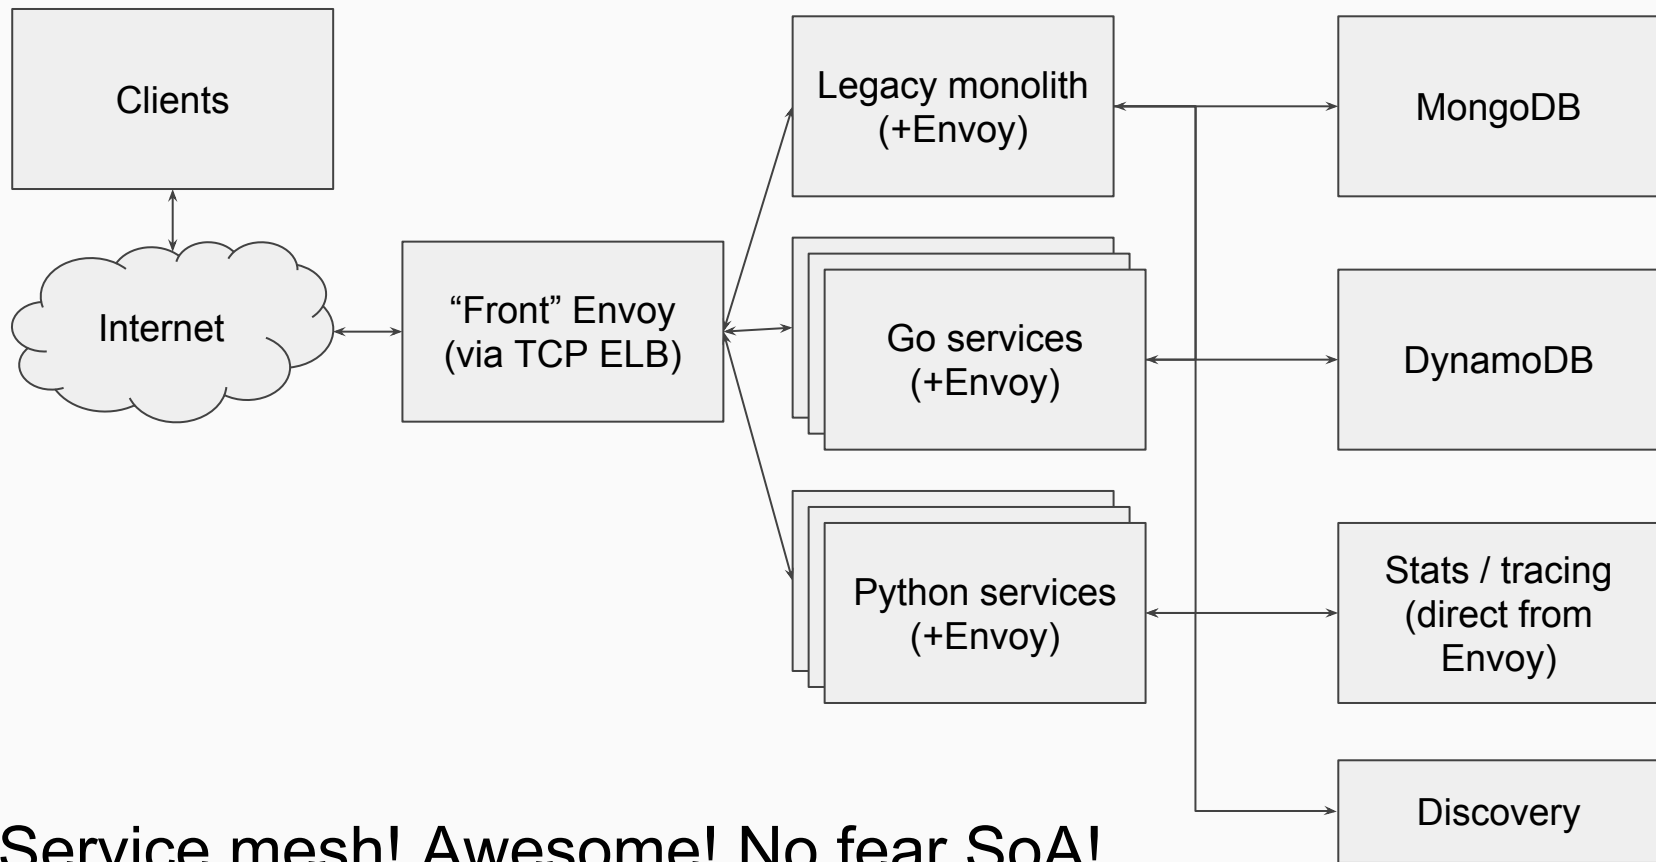
What is Envoy

- **Out of process architecture:** Let's do a lot of really hard stuff in one place and allow application developers to focus on business logic.
- **Modern C++11 code base:** Fast and productive.
- **L3/L4 filter architecture:** A byte proxy at its core. Can be used for things other than HTTP (e.g., MongoDB, redis, stunnel replacement, TCP rate limiter, etc.).
- **HTTP L7 filter architecture:** Make it easy to plug in different functionality.
- **HTTP/2 first!** (Including **gRPC** and a nifty gRPC HTTP/1.1 bridge).
- **Service discovery and active/passive health checking.**
- **Advanced load balancing:** Retry, timeouts, circuit breaking, rate limiting, shadowing, outlier detection, etc.
- Best in class **observability:** stats, logging, and tracing.
- **Edge proxy:** routing and TLS.

Envoy service to service topology



Lyft today



Service mesh! Awesome! No fear SoA!

Eventually consistent service discovery

- **Fully consistent** service discovery systems are very popular (ZK, etcd, consul, etc.).
- In practice they are **hard to run at scale**.
- Service discovery is actually an **eventually consistent** problem. Let's recognize that and design for it.
- Envoy is designed from the get go to treat **service discovery as lossy**.
- Active health checking used in combination with service discovery to produce a **routable overlay**.

Discovery Status	HC OK	HC Failed
Discovered	Route	Don't Route
Absent	Route	Don't Route / Delete

Advanced load balancing

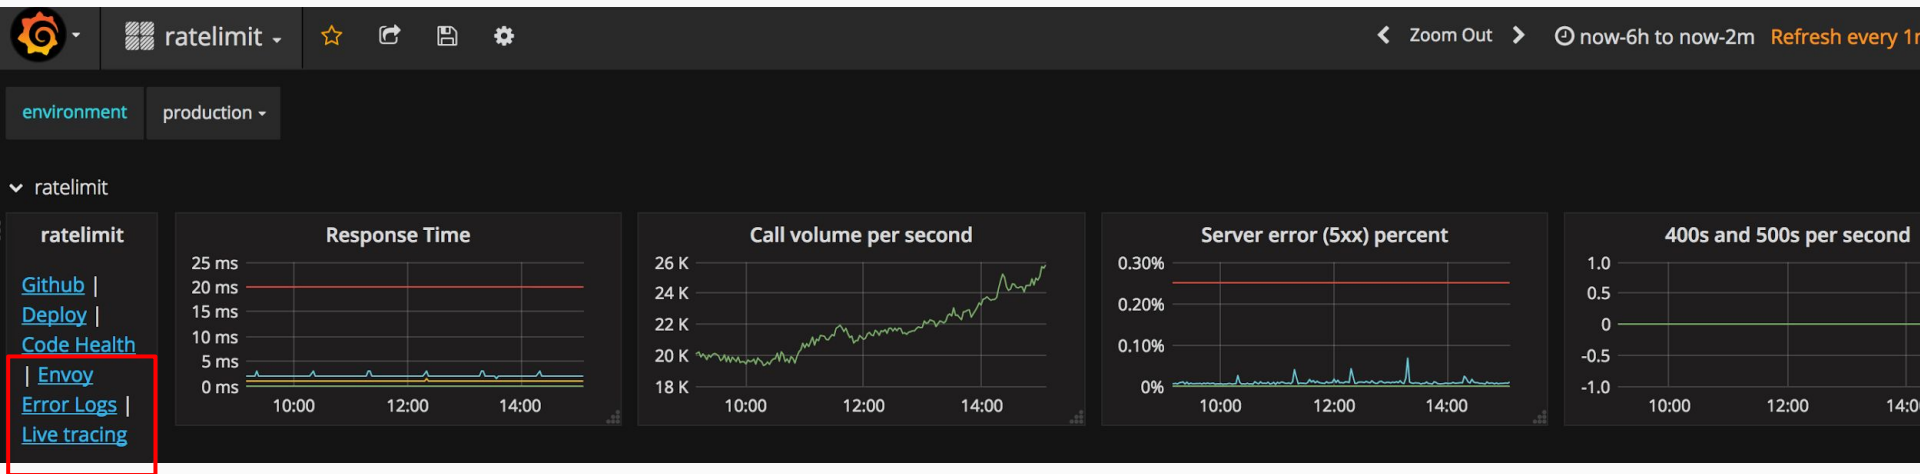
- Different **service discovery** types.
- Zone aware least request **load balancing**.
- **Dynamic stats**: Per zone, canary specific stats, etc.
- **Circuit breaking**: Max connections, requests, and retries.
- **Rate limiting**: Integration with global rate limit service.
- **Shadowing**: Fork traffic to a test cluster.
- **Retries**: HTTP router has built in retry capability with different policies.
- **Timeouts**: Both “outer” (including all retries) and “inner” (per try) timeouts.
- **Outlier detection**: Consecutive 5xx
- **Deploy control**: Blue/green, canary, etc.
- **Fault injection**

Observability

- **Observability** is by far the most important thing that Envoy provides.
- Having all SoA traffic transit through Envoy gives us a single place where we can:
 - Produce consistent **statistics** for every hop
 - Create and propagate a stable **request ID / tracing context**
 - Consistent **logging**
 - Distributed **tracing**




Observability: Per service auto-generated panel



Links to logging and tracing

Observability: Service to service template dashboard

Template with drop down for every service



Observability: Envoy global health dashboard

Envoy-Global








Zoom Out

6 hours ago to 2 minutes ago refreshed ev

DASHBOARD README

TOP LEVEL ALL ENVOYS




CROSS ZONE TRAFFIC

RATELIMIT

Observability: Distributed tracing

Expand all spans



Observability: Logging



Discover

Visualize

Dashboard

Settings

unique_id:82515ede-9a51-9c63-b47d-002a6cf74471

[logstash-]YYYY.MM.DD-HH

Selected Fields

? _source

Available Fields



Popular

duration_millis

† http_version

† source

Quick Count ⓘ (6 / 6 records)

/var/log/envoy/ingress_http_err... 🔍

50.0%

/var/log/envoy/access_error.log 🔍

50.0%

Visualize

status

@size

October 30th 2016, 20:38:26.391 - October 30th 2016, 20:53:26.391 — [by 30 seconds](#)



Time ▾

_source

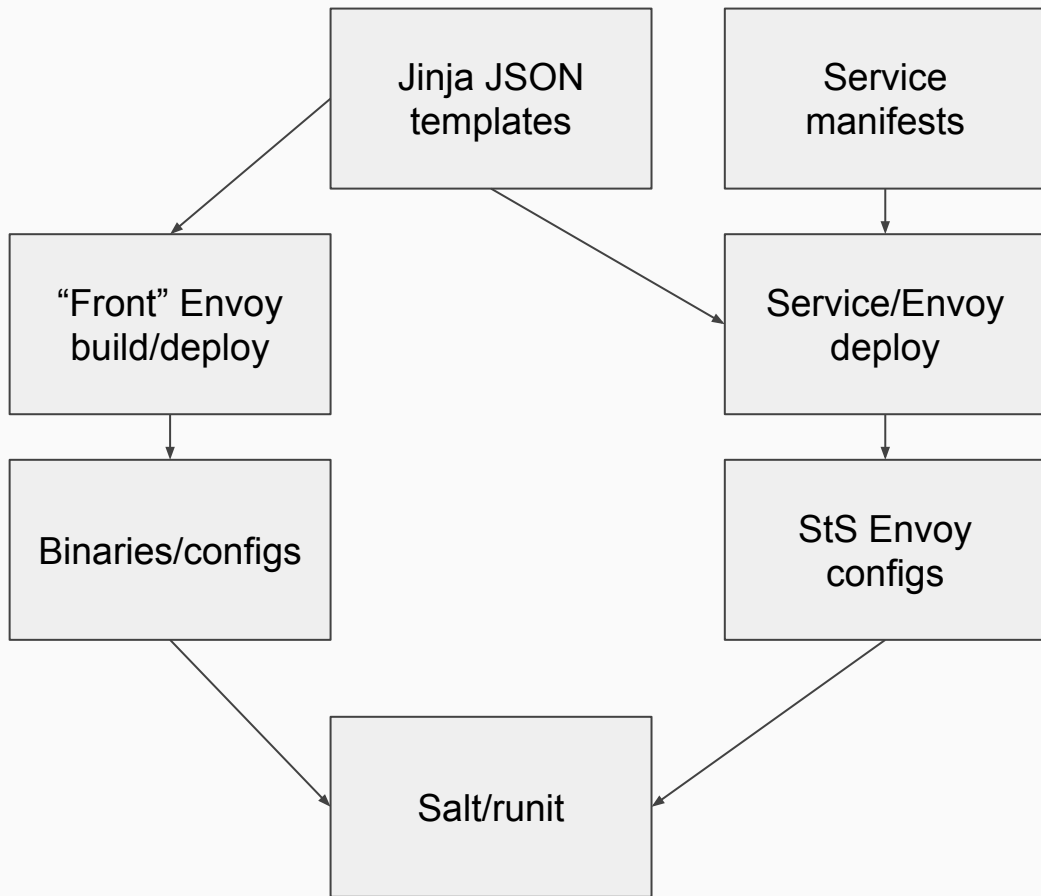
▶ October 30th 2016, 20:51:38.001 **unique_id:** 82515ede-9a51-9c63-b47d-002a6cf74471 **asg:** envoy **bytes_received:** 2 **bytes_sent:** 0
0.5.22 **duration_millis:** 556 **host:** envoy-production-iad-8c37a31f **host_header:** api-internal.lyft
P/1.1 **method:** POST **search.uid:** 82515ede-9a51-9c63-b47d-002a6cf74471 **search.lyft_id:** 77392617
log/envoy/access_error.log **status:** 500 **upstream_ip:** tcp://10.0.127.39:9211 **upstream_time_milli**
s/773926173699609484/driverStatus **user_agent:** jobscheduler:worker:0.1:1.0 **@timestamp:** October :

Performance matters for a service proxy

- For most companies **developer time is worth more than infra costs** (cost vs. throughput).
- However, **Latency** and **predictability** is what matters. And in particular **tail latency** (P99+).
- Virtual IaaS, multiple languages and runtimes, languages that use GC: Niceties that improve productivity and reduce upfront dev costs, but make **debugging really difficult**.
- Ability to **reason about overall performance** and reliability is critical.



Envoy config/process management @Lyft



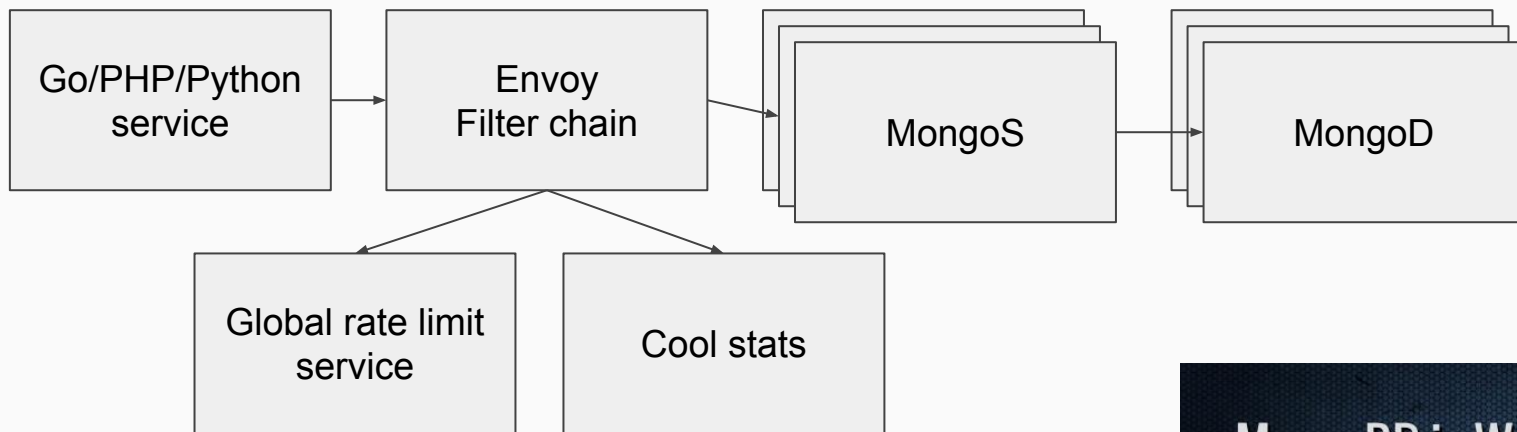
- Combination of static and dynamic configs.
- Service egress, circuit breaking, etc. configs specified in manifest.
- Service configs built on service host at service/envoy deploy time.
- **Next up for Lyft: config service via APIs!**

Envoy thin clients @Lyft

```
from lyft.api_client import EnvoyClient
switchboard_client = EnvoyClient(
    service='switchboard'
)
msg = {'template': 'breaksignout'}
headers = {'x-lyft-user-id': 12345647363394}
switchboard_client.post("/v2/messages", data=msg, headers=headers)
```

- Abstract away egress port
- Request ID/tracing propagation
- Guide devs into good timeout, retry, etc. policies
- Similar thin clients for Go and PHP

Envoy MongoDB proxy @Lyft



No more death spirals! Web scale!

Filters:

- L4 global rate limit (limit CPS into MongoS)
- Mongo sniffer (cool stats)
- TCP proxy (MongoS load balancing)


MongoDB is Web Scale



Envoy deployment @Lyft

- > 100 services.
- > 10,000 hosts.
- > 2,000,000 RPS.
- All service to service traffic (REST and gRPC).
- Use gRPC bridge to unlock Python gevent clients.
- MongoDB proxy.
- DynamoDB proxy.
- External service proxy (AWS and other partners).
- Kibana/Elastic Search for logging.
- LightStep for tracing.
- Wavefront for stats (via statsd).

- Thanks for coming!
- We are super excited about building a **community** around Envoy. Talk to us if you need help getting started.
- <https://lyft.github.io/envoy/>
- **Lyft is hiring**: Contact us if you want to work on hard scaling problems in a fast moving company: <https://www.lyft.com/jobs>



Envoy edge proxy topology

